Managing and Manipulating Data Using R

Lecture 1.1

Karina Salazar

1. Student introductions
2. About your instructor
3. What is R
4. What is this course about?
5. Course logistics
6. Create “R project” and directory structure
7. Directories and filepaths

1 Student introductions

Student introductions

1. Preferred name
2. Preferred pronouns
3. Academic program (and how far along)
4. GA, RA, TA, and/or job?
5. Why are you interested in this course?
6. Tell us about your name (e.g., what is the origin story?)

2 About your instructor

Karina Salazar, instructor

My start in data management/statistical analysis

- SPSS
  - evaluated retention programs within institutional research and assessment offices
  - student-level data on math remediation courses
  - College Academy for Parents, Think Tank, Assessment Institute
- Stata
  - used loops and user-defined functions to work with national datasets (IPEDS, Survey of Earned Doctorates)

Got sick of the limitations of survey data and/or available data

- No survey asked questions on what I was interested in
  - universities pledge commitment to access, but enrollments don't tell the whole story
  - who do they actually recruit?
- We realized "data science" could create data from publicly available data sources
  - Twitter
  - travel schedules on admissions websites

Recruiting research program and “data science”

- Python
  - web-scraping
  - connecting to Application Program Interfaces (API) (e.g., census data, Twitter, LinkedIn)
  - Natural Language Processing
- R
  - R can do all “data science” tasks Python can
  - R can do all statistical analyses that Stata can (and more!)
  - R has amazing mapping capabilities

Examples:

- The off-campus recruiting project
- Dissertation Defense

3 What is R

What is R

According to the Inter-university consortium for political and social research (ICPSR):

> *R is "an alternative to traditional statistical packages such as SPSS, SAS, and Stata such that it is an extensible, open-source language and computing environment for Windows, Macintosh, UNIX, and Linux platforms. Such software allows for the user to freely distribute, study, change, and improve the software under the Free Software Foundation's GNU General Public License."*

- For more info visit R-project.org

Base R vs. R packages

There are "default" packages that come with R. Some of these include:

- `as.character`
- `print`
- `setwd`

And there are R packages developed and shared by others. Some R packages include:

- `tidyverse`
- `stargazer`
- `foreign`

more about these in later weeks...

Installing and Loading R packages

You only need to install a package once. To install an R package use `install.package()` function.

```
#install.packages("tidyverse")
```

However, you need to load a package everytime you plan to use it. To load a package use the `library()` function.

```
library(tidyverse)
#> -- Attaching packages ---------------------------------------------
#> v ggplot2 3.2.1     v purrr   0.3.2
#> v tibble  2.1.3     v dplyr   0.8.3
#> v tidyr   0.8.3     v stringr 1.4.0
#> v readr   1.3.1     v forcats 0.4.0
#> -- Conflicts ------------------------------------------------------
#> x dplyr::filter() masks stats::filter()
#> x dplyr::lag()    masks stats::lag()
```

RStudio

"RStudio is an integrated development environment (IDE) for R. It includes a console, syntax-highlighting editor that supports direct code execution, as well as tools for plotting, history, debugging and workspace management."

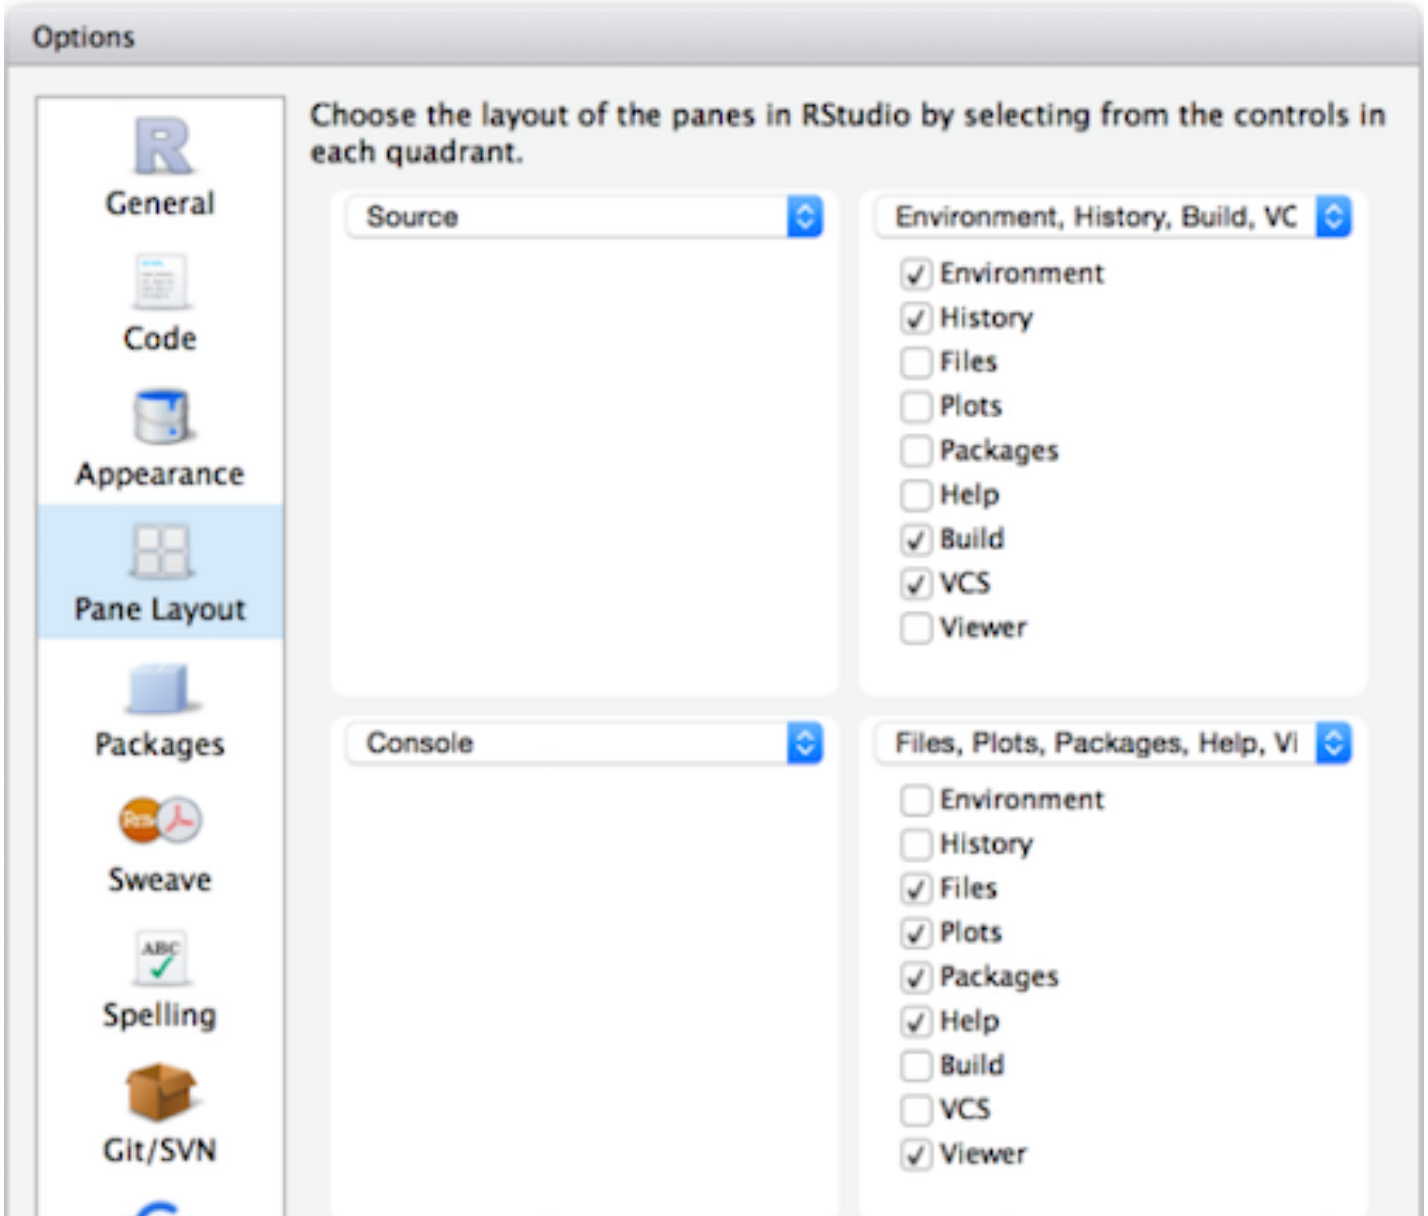

R Markdown

R Markdown produces dynamic output formats in html, pdf, MS Word, dashboards, Beamer presentations, etc.

- We will be using R Markdown for lectures and homeworks.

Why R? Capabilities of R

- Graphs
- Presentation
- Websites
- Journals
- Interactive tutorials
- Web apps
- Dashbaords
- Books
- Web scraping
- Maps

For more info visit

Graphs

- Create graphs with ggplot2 package

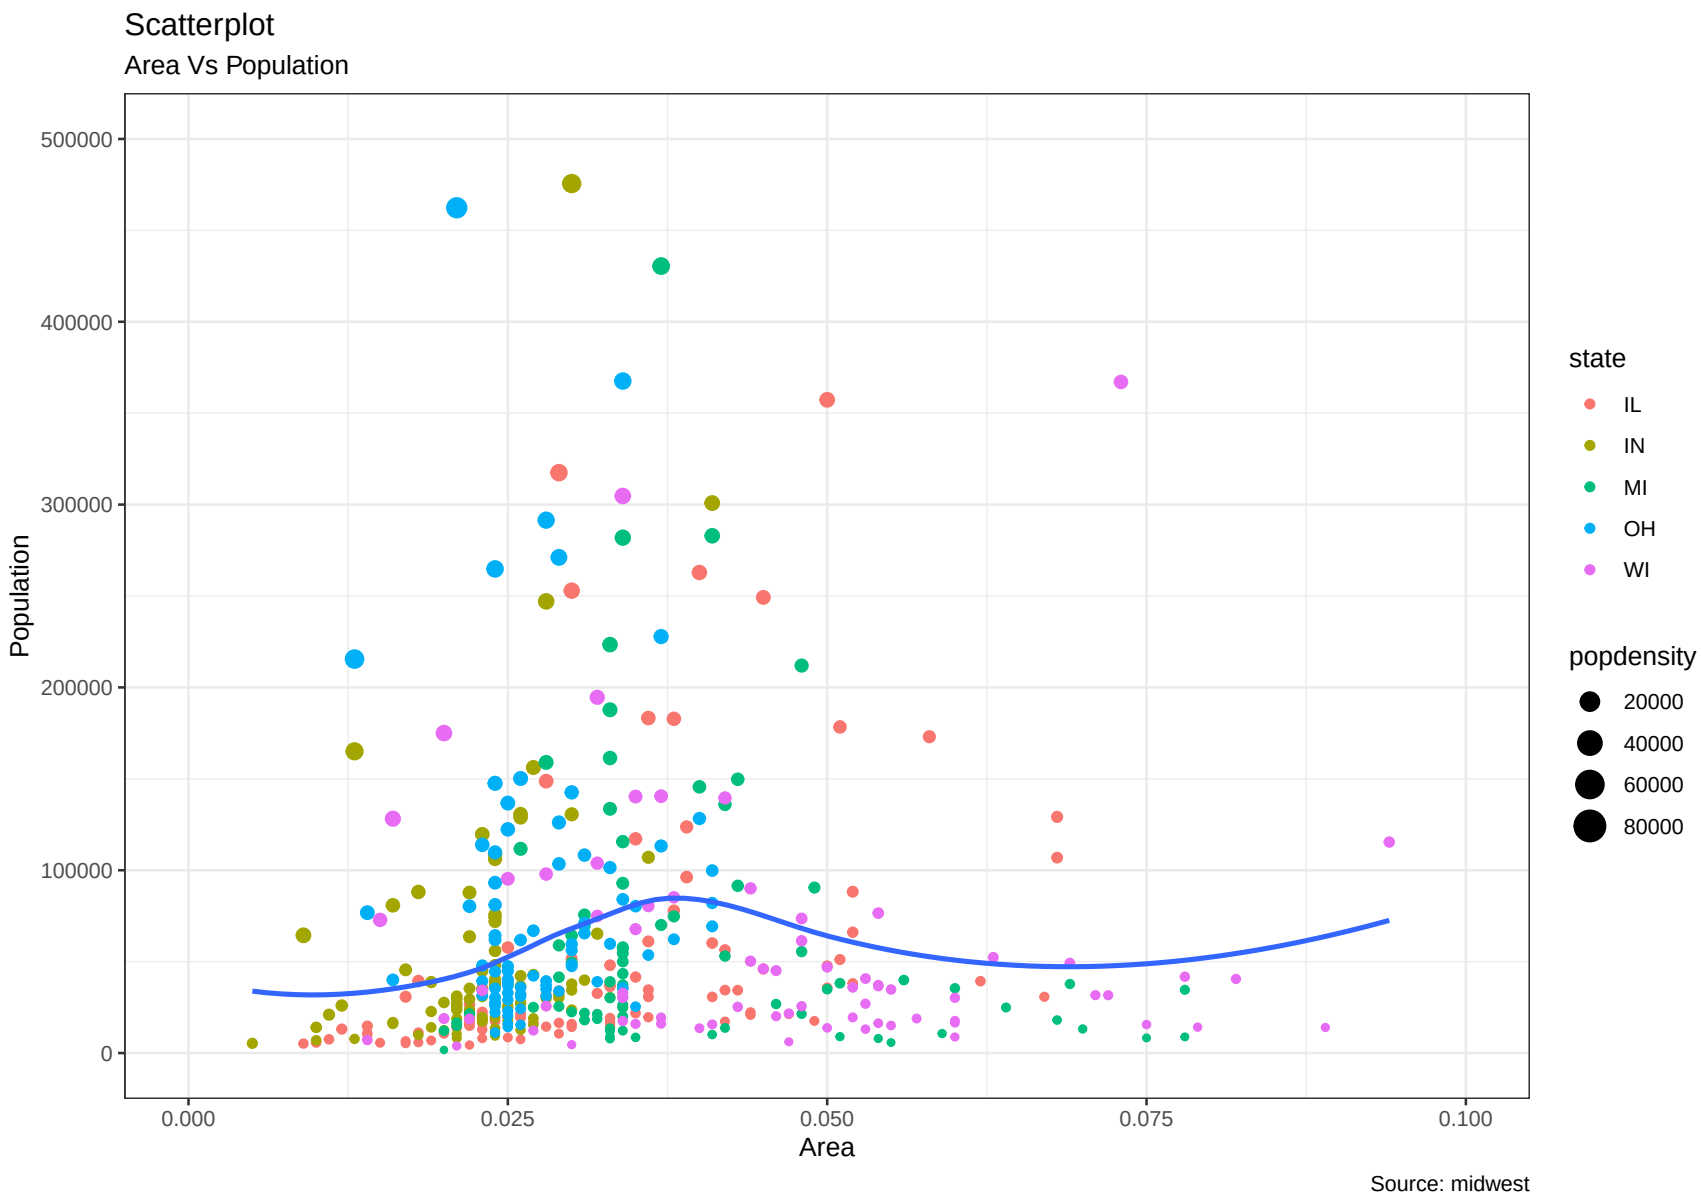

Journal articles

- Journal articles with rticles package

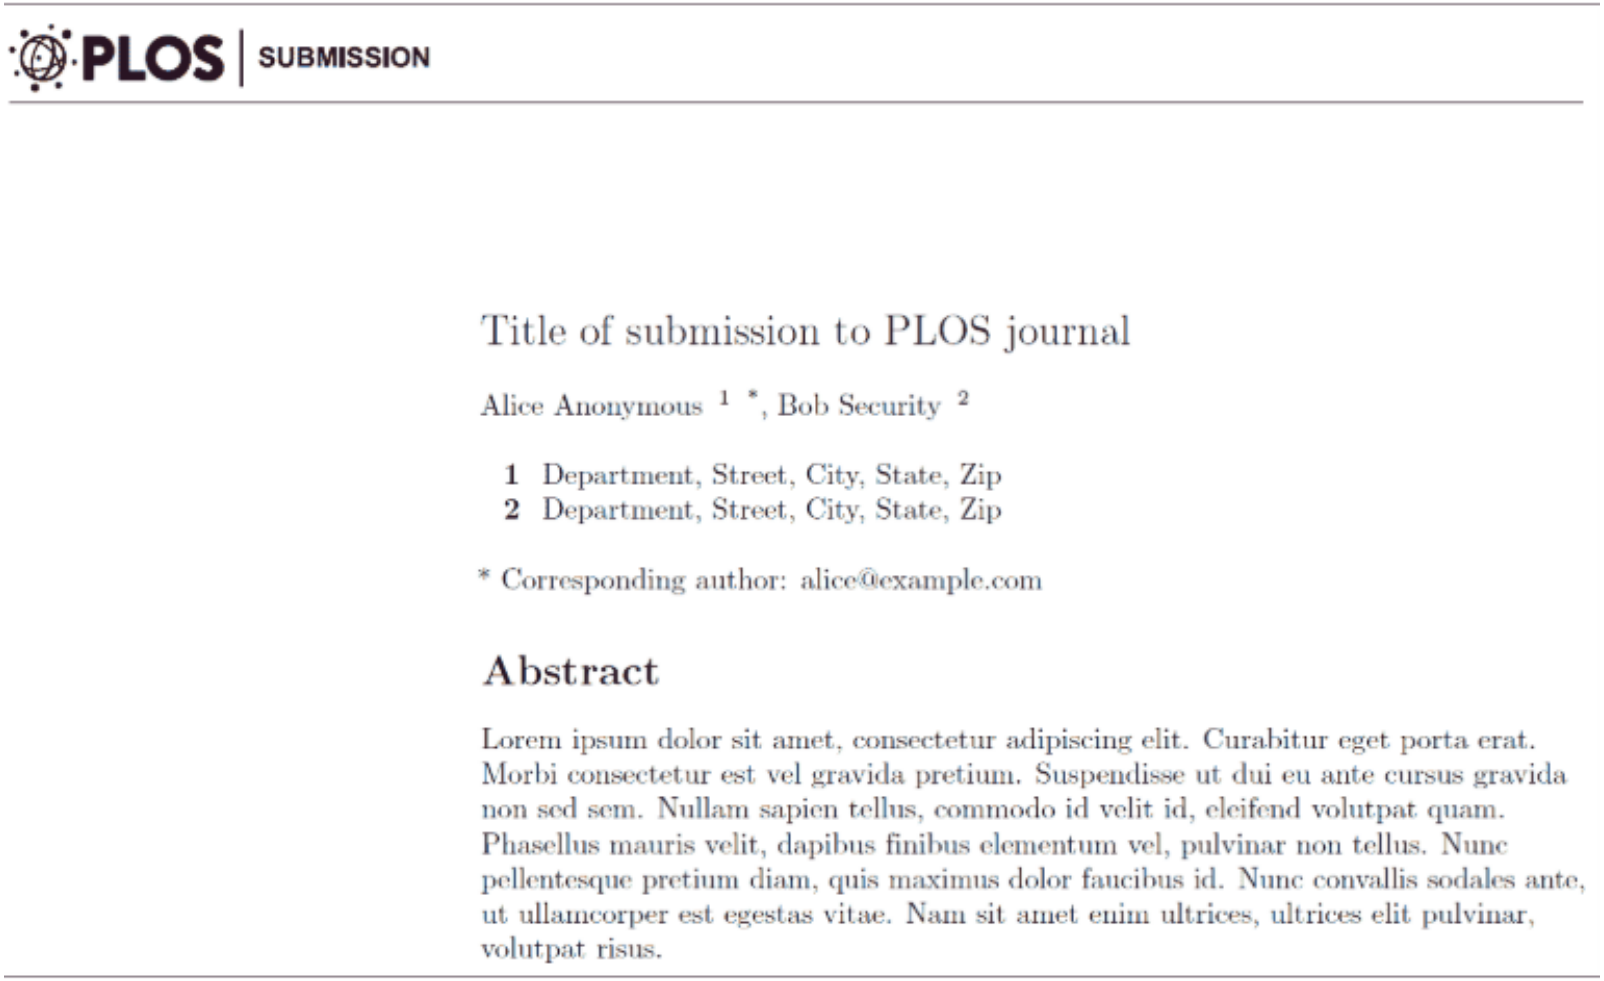

PLOS | SUBMISSION

Title of submission to PLOS journal

Alice Anonymous [1] *, Bob Security [2]

1 Department, Street, City, State, Zip
2 Department, Street, City, State, Zip

* Corresponding author: alice@example.com

Abstract

Lorem ipsum dolor sit amet, consectetur adipiscing elit. Curabitur eget porta erat. Morbi consectetur est vel gravida pretium. Suspendisse ut dui eu ante cursus gravida non sed sem. Nullam sapien tellus, commodo id velit id, eleifend volutpat quam. Phasellus mauris velit, dapibus finibus elementum vel, pulvinar non tellus. Nunc pellentesque pretium diam, quis maximus dolor faucibus id. Nunc convallis sodales ante, ut ullamcorper est egestas vitae. Nam sit amet enim ultrices, ultrices elit pulvinar, volutpat risus.

Interactive web apps

- Interactive web apps with shiny package

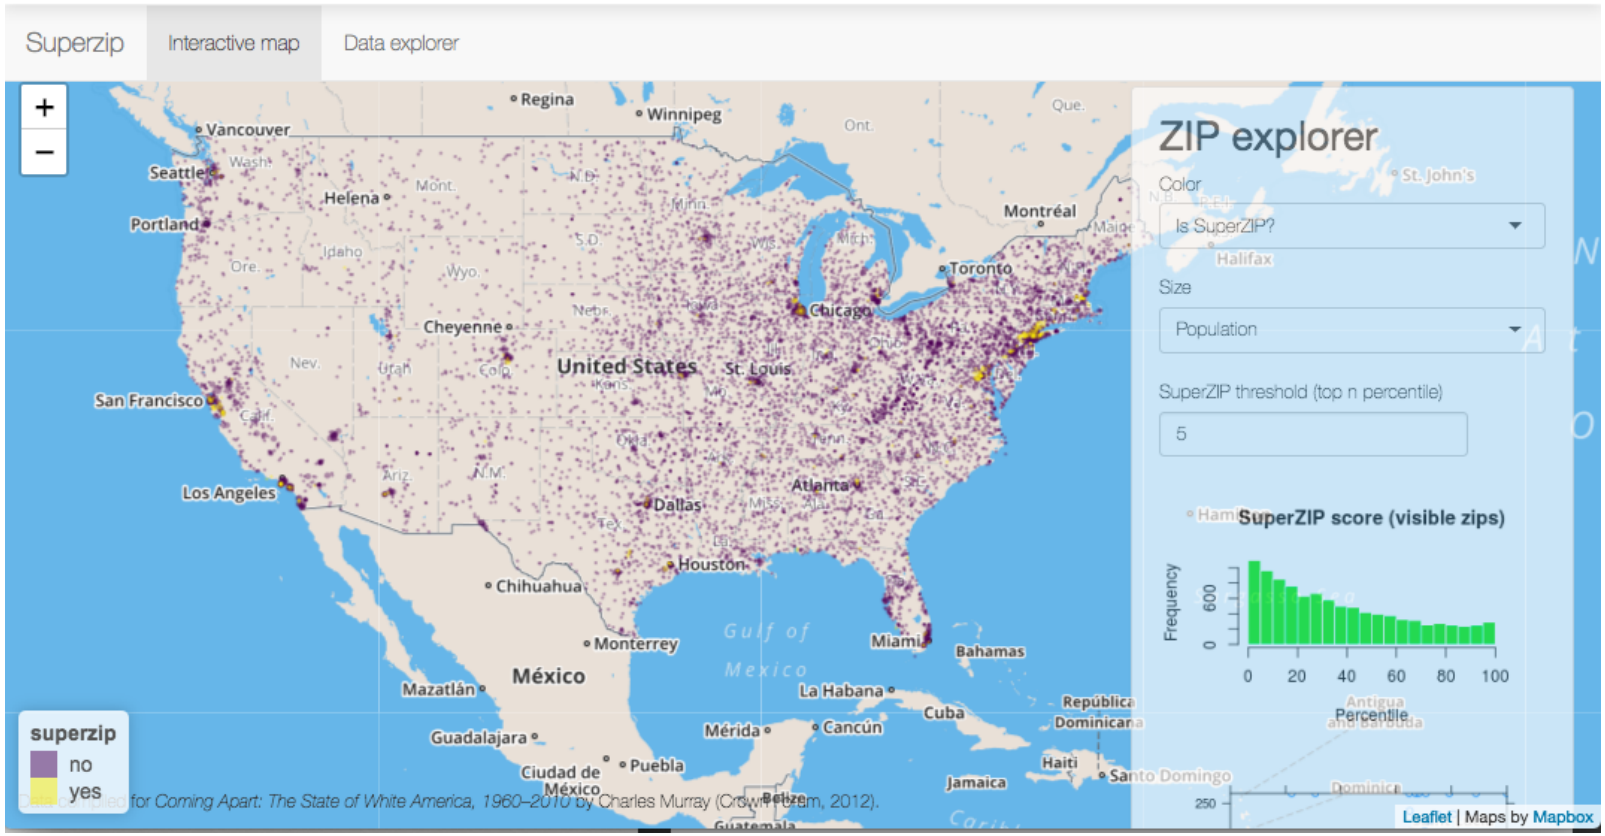

Mapping

- Mapping with sf package & ggplot

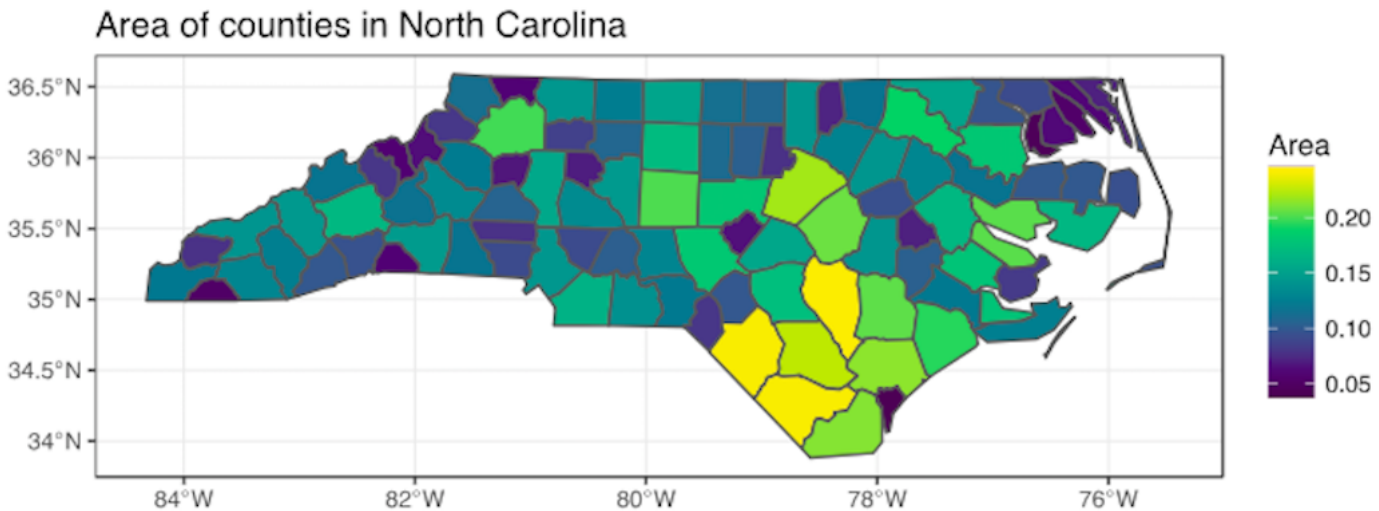

4 What is this course about?

What is data management?

- All the stuff you have to do to create analysis datasets that are ready to analyze:
  - collect data
  - read/import data into statistical programming language
  - clean data
  - integrate data from multiple sources (e.g, join/merge, append)
  - change organizational structure of data so it is suitable for analysis
  - create “analysis variables” from “input variables”
  - Make sure that you have created analysis variables correctly

Why I don't call this class "R for data science"

Learn to walk before you can run!

- "data science" implies doing "fancy" things like mapping, network analysis, web-scraping, etc.
- But if you don't know how to clean data, these "fancy" analyses and visualizations will be useless
- "80% of data science is data cleaning"
- The skills you learn in this data management class are foundational to data science tasks! (and a prerequisite to taking data science seminar)

Who is this class for?

This class is for anyone who wants to work with data, that is people who want to be:

- researchers working with survey data and doing traditional statistical analyses
- researchers who want to do “data science” oriented research involving
- analysts working at think tanks or non-profits
- “Data scientists”

5 Course logistics

Course logistics

- follow the syllabus

6 Create “R project” and directory structure

What is an R project? Why are you doing this?

What is an “R project”?

- helps you keep all files for a project in one place
- When you open an R project, the file-path of your current working directory is automatically set to the file-path of your R-project

Why am I asking you to create R project and download a specific directory structure?

- I want you to be able to run the .Rmd files for each lecture on your own computer
- Sometimes these .Rmd files point to certain sub-folders
- If you create R project and create directory structure I recommend, you will be able to run .Rmd files from your own computer without making any changes to file-paths!

Follow these steps to create "R project" and directory structure

1. Download this zip folder: LINK HERE
   - Unzip the folder: this is a shell of the file directory you should use for this class
   - Change the name to "rclass"
   - Move it to your preferred location (e.g, documents, desktop, dropbox, etc)
2. In RStudio, click on "File" >> "New Project" >> "New Directory" >> "New Project"
   - In "Directory name", type "rclass_project" as the title of the Rproject for the course
   - In "Create project as subdirectory of", click browse and:
     - save the R Project within the rclass folder (same general folder as data and lectures)
3. Save the following files in "rclass/lectures/lecture1"
   - lecture1.1_ua.Rmd
   - lecture1.1_ua.pdf
   - lecture1.2_ua.Rmd
   - lecture1.2_ua.Pdf
   - lecture1.2_ua.R

After you follow these steps

- you can add any additional sub-folders you want to the “rclass” folder
  - e.g., “syllabus”, “resources”
- You can add any additional files you want to the sub-directory folders you unzipped
  - e.g., in “rclass/lectures/lecture1” you might add an additional document of notes you took

7 Directories and filepaths

Working directory

(Current) Working directory

- the folder/directory in which you are currently working
- this is where R looks for files
- Files located in your current working directory can be accessed without specifying a filepath because R automatically looks in this folder

Function `getwd()` shows current working directory

```
getwd()
#> [1] "/Users/Karina/rclass/lectures/lecture1"
```

Command `list.files()` lists all files located in working directory

```
getwd()
#> [1] "/Users/Karina/rclass/lectures/lecture1"
list.files()
#>  [1] "data-structures-overview.png" "images.zip"
#>  [3] "lecture1.1_files"             "lecture1.1_ua_files"
#>  [5] "lecture1.1_ua.pdf"            "lecture1.1_ua.Rmd"
#>  [7] "lecture1.1_ua.tex"            "lecture1.1.pdf"
#>  [9] "lecture1.1.Rmd"               "lecture1.2_ua.pdf"
#> [11] "lecture1.2_ua.Rmd"            "lecture1.2.pdf"
#> [13] "lecture1.2.R"                 "lecture1.2.Rmd"
#> [15] "lecture1.pdf"                 "lecture1.Rmd"
#> [17] "pane_layout.png"              "problemset1_solutions.pdf"
#> [19] "problemset1_solutions.Rmd"    "problemset1.pdf"
#> [21] "problemset1.Rmd"              "rticles.png"
#> [23] "sample_hw_oi.html"            "sample_hw_oi.Rmd"
```

Working directory, "Code chunks" vs. "console" and "R scripts"

When you run **code chunks** in RMarkdown files (.Rmd), the working directory is set to the filepath where the .Rmd file is stored

```
getwd()
#> [1] "/Users/Karina/rclass/lectures/lecture1"
list.files()
#>  [1] "data-structures-overview.png" "images.zip"
#>  [3] "lecture1.1_files"             "lecture1.1_ua_files"
#>  [5] "lecture1.1_ua.pdf"            "lecture1.1_ua.Rmd"
#>  [7] "lecture1.1_ua.tex"            "lecture1.1.pdf"
#>  [9] "lecture1.1.Rmd"               "lecture1.2_ua.pdf"
#> [11] "lecture1.2_ua.Rmd"            "lecture1.2.pdf"
#> [13] "lecture1.2.R"                 "lecture1.2.Rmd"
#> [15] "lecture1.pdf"                 "lecture1.Rmd"
#> [17] "pane_layout.png"              "problemset1_solutions.pdf"
#> [19] "problemset1_solutions.Rmd"    "problemset1.pdf"
#> [21] "problemset1.Rmd"              "rticles.png"
#> [23] "sample_hw_oj.html"            "sample_hw_oj.Rmd"
#> [25] "sf.png"                       "shiny.png"
```

When you run code from the **R Console** or an **R Script**, the working directory is....

Command `getwd()` shows current working directory

```
getwd()
#> [1] "/Users/Karina/rclass/lectures/lecture1"
```

Absolute vs. relative filepath

Absolute file path: The absolute file path is the complete list of directories needed to locate a file or folder.

```
setwd("Users/Karina/rclass/lectures/lecture2")
```

Relative file path: The relative file path is the path relative to your current location/directory. Assuming your current working directory is in the "lecture2" folder and you want to change your directory to the data folder, your relative file path would look something like this:

```
setwd("../../data")
```

File path shortcuts

Key	Description
~	tilde is a shortcut for user's home directory (mine is my name)
../	moves up a level
../../	moves up two level