

Managing and Manipulating Data Using R

Lecture 4

Ozan Jaquette

1. Introduction

2. `group_by()` and `summarise()`

Introduction

Libraries we will use today

```
library(tidyverse)
```

```
#> -- Attaching packages -----
```

```
#> v ggplot2 3.0.0      v purrr  0.2.5
```

```
#> v tibble  1.4.2      v dplyr  0.7.6
```

```
#> v tidyr   0.8.1      v stringr 1.3.1
```

```
#> v readr   1.1.1      v forcats 0.3.0
```

```
#> -- Conflicts -----
```

```
#> x dplyr::filter() masks stats::filter()
```

```
#> x dplyr::lag()    masks stats::lag()
```

tidyverse

Data we will use today

Data on off-campus recruiting events by public universities

- Object `df_event`
 - One observation per university, recruiting event
- Object `df_event`
 - One observation per high school (visited and non-visited)

```
rm(list = ls()) # remove all objects

#load dataset with one obs per recruiting event
load("../..data/recruiting/recruit_event_somevars.Rdata")

#load dataset with one obs per high school
load("../..data/recruiting/recruit_school_somevars.Rdata")

load("../..data/prospect_list/western_washington_college_board_list.RData")
```

Processing across observations, introduction

Creation of analysis datasets often requires calculations across obs

Examples:

- You have a dataset with one observation per student-term and want to create a variable of credits attempted per term
- You have a dataset with one observation per student-term and want to create a variable of GPA for the semester or cumulative GPA for all semesters
- Number of off-campus recruiting events university makes to each state
- Average household income at visited versus non-visited high schools

Processing across variables vs. processing across observations

Visits by UC Berkely to public high schools

```
#> # A tibble: 5 x 6
#>   school_id    state tot_stu_pub fr_lunch pct_fr_lunch med_inc
#>   <chr>        <chr>    <dbl>    <dbl>      <dbl>    <dbl>
#> 1 340882002126 NJ          1846      29      0.0157 178732
#> 2 340147000250 NJ          1044      50      0.0479  62288
#> 3 340561003796 NJ          1505     298      0.198 100684.
#> 4 340165005124 NJ          1900      43      0.0226 160476.
#> 5 341341003182 NJ          1519     130      0.0856 144346
```

- So far, we have focused on “processing across variables”
 - ▷ Performing calculations across columns (i.e., vars), typically within a row (i.e., observation)
 - ▷ Example: percent free-reduced lunch
- Processing across obs (focus of today’s lecture)
 - ▷ Performing calculations across rows (i.e., obs), often within a column (i.e., variable)
 - ▷ Example: Average household income of visited high schools, by state

`group_by()` and `summarise()`

group_by()

`group_by()` converts a data frame object into groups. After grouping, functions performed on data frame are performed “by group”

- part of **dplyr** package within **tidyverse**; not part of **Base R**
- works best with pipes `%>%` and `summarise()` function [described below]

Basic syntax:

- `group_by(object, vars to group by separated by commas)`

Typically, “group_by” variables are character, factor, or integer variables

Possible “group by” variables in `df_event` data

- university
- event type (e.g., public HS, private HS, hotel)
- state

group_by()

Group `df_event` data by university, event type, and event state

- `group_by` doesn't do much by itself; just prints data

```
group_by(df_event, univ_id, event_type, event_state)
```

```
df_event %>% group_by(univ_id, event_type, event_state) # using pipes
```

Grouping is not retained unless you **assign** it

```
class(df_event)
#> [1] "tbl_df"      "tbl"        "data.frame"
df_event_grp <- df_event %>% group_by(univ_id, event_type, event_state) # using
class(df_event_grp)
#> [1] "grouped_df" "tbl_df"     "tbl"        "data.frame"
```

Use `ungroup(object)` to un-group grouped data

```
df_event_grp <- ungroup(df_event_grp)
class(df_event_grp)
#> [1] "tbl_df"      "tbl"        "data.frame"
rm(df_event_grp)
```

summarise()

`summarise()` function performs calculations across rows of a data frame and then collapses the data frame to a single row

Basic syntax [see documentation]:

- `summarise(object, summarise functions separated by commas)`
- summarise functions include: `n()`, `mean()`, `first()`, etc.

Simple example (output omitted)

```
summarise(df_event, num_events=n())  
df_event %>% summarise(num_events=n()) # using pipes
```

Object created by `summarise()` not retained unless you **assign** it

```
event_temp <- df_event %>% summarise(num_events=n(),  
  mean_inc=mean(med_inc, na.rm = TRUE))  
  
event_temp  
#> # A tibble: 1 x 2  
#>   num_events mean_inc  
#>   <int>     <dbl>  
#> 1    17976    88774.  
rm(event_temp)
```

I'll explain `na.rm = TRUE` later

Combining summarise() and group_by

summarise() on ungrouped vs. grouped data:

- By itself, summarise() performs calculations across all rows of data frame then collapses the data frame to a single row
- When data frame is grouped, summarise() performs calculations across rows within a group and then collapses to a single row for each group

Number of recruiting events for each university

```
df_event %>% group_by(instnm) %>% summarise(num_events=n())
```

Number of recruiting events by event_type for each university

```
df_event %>% group_by(instnm, event_type) %>% summarise(num_events=n())
```

Number of events and avg. pct White by event_type for each university

```
df_event %>% group_by(instnm, event_type) %>%  
  summarise(num_events=n(),  
            mean_pct_white=mean(pct_white_zip, na.rm = TRUE)  
  )
```

Combining summarise() and group_by

Recruiting events by UC Berkeley

```
df_event %>% filter(univ_id == 110635) %>%  
  group_by(event_type) %>% summarise(num_events=n())
```

Let's create a dataset of recruiting events at UC Berkeley

The 0/1 variable event_inst equals 1 if event is in same state as the university

```
#event_berk %>% group_by(event_type, event_inst) %>% select(pid, event_date, event_type, event_inst)  
event_berk %>% arrange(event_date) %>% select(pid, event_date, event_type, event_inst)  
#> # A tibble: 8 x 5  
#>   pid event_date event_type event_state event_inst  
#>   <int> <date>      <fct>      <chr>      <chr>  
#> 1 13100 2017-04-11 other      HI          Out-State  
#> 2 13089 2017-04-14 public hs   GA          Out-State  
#> 3 13088 2017-04-23 private hs CT          Out-State  
#> 4 13086 2017-04-23 other      CA          In-State  
#> 5 13091 2017-04-24 private hs NY          Out-State  
#> 6 13087 2017-04-24 public hs   CA          In-State  
#> 7 13092 2017-04-25 other      NY          Out-State  
#> 8 13099 2017-04-25 2yr college CA          In-State
```

summarise(): Counts

The count function `n()` takes no arguments and returns the size of the current group

```
event_berk %>% group_by(event_type, event_inst) %>%  
  summarise(num_events=n())
```

Object not retained unless we **assign**

```
berk_temp <- event_berk %>% group_by(event_type, event_inst) %>%  
  summarise(num_events=n())  
berk_temp  
typeof(berk_temp)  
str(berk_temp)
```

Because counts are so important, `dplyr` package includes separate `count()` function that can be called outside `summarise()` function

```
event_berk %>% group_by(event_type, event_inst) %>% count()  
event_berk %>% group_by(event_type) %>% count(event_inst) # same  
  
berk_temp2 <- event_berk %>% group_by(event_type, event_inst) %>% count()  
berk_temp == berk_temp2  
rm(berk_temp,berk_temp2)
```

summarise(): count with logical vectors and sum()

Logical vectors have values TRUE and FALSE.

- When used with numeric functions, TRUE converted to 1 and FALSE to 0.

sum() is a numeric function that returns the sum of values

```
sum(c(5,10))  
#> [1] 15  
sum(c(TRUE,TRUE,FALSE,FALSE))  
#> [1] 2
```

is.na() returns TRUE if value is NA and otherwise returns FALSE

```
is.na(c(5,NA,4,NA))  
#> [1] FALSE TRUE FALSE TRUE
```

Application: How many missing/non-missing obs in variable **[very important]**

```
event_berk %>% group_by(event_type) %>%  
  summarise(  
    n_events = n(),  
    n_miss_inc = sum(is.na(med_inc)),  
    n_nonmiss_inc = sum(!is.na(med_inc)),  
    n_nonmiss_fr_lunch = sum(!is.na(fr_lunch))  
  )
```


summarise(): means

The `mean()` function within `summarise()` calculates means, separately for each group

```
event_berk %>% group_by(event_inst, event_type) %>% summarise(
  n_events=n(),
  mean_inc=mean(med_inc, na.rm = TRUE),
  mean_pct_white=mean(pct_white_zip, na.rm = TRUE)) %>% head(5)
```

#> # A tibble: 5 x 5

#> # Groups: event_inst [1]

#>	event_inst	event_type	n_events	mean_inc	mean_pct_white
#>	<chr>	<fct>	<int>	<dbl>	<dbl>
#> 1	In-State	public hs	260	87146.	39.8
#> 2	In-State	private hs	36	94133.	48.0
#> 3	In-State	2yr college	107	79144.	40.5
#> 4	In-State	4yr college	12	148587.	55.1
#> 5	In-State	other	50	73218.	35.9

I'll talk about `na.rm = TRUE` on next slide

summarise(): means and NA values

The default behavior of “aggregation functions” (e.g., `summarise()`) is if the **input** has any missing value (NA) than the output will be missing.

`na.rm` (in words “remove NA”) is an option available in many functions.

- `na.rm = FALSE` [the default for `mean()`]
 - ▷ Do not remove missing values from input before calculating
 - ▷ Therefore, missing values in input will cause output to be missing
- `na.rm = TRUE`
 - ▷ Remove missing values from input before calculating
 - ▷ Therefore, missing values in input will not cause output to be missing

#na.rm = FALSE; the default setting

```
event_berk %>% group_by(event_inst, event_type) %>% summarise(
  n_events=n(),
  n_miss_inc = sum(is.na(med_inc)),
  mean_inc=mean(med_inc, na.rm = FALSE),
  n_miss_frlunch = sum(is.na(fr_lunch)),
  mean_fr_lunch=mean(fr_lunch, na.rm = FALSE))
```

#na.rm = TRUE

```
event_berk %>% group_by(event_inst, event_type) %>% summarise(
  n_events=n(),
  n_miss_inc = sum(is.na(med_inc)),
  mean_inc=mean(med_inc, na.rm = TRUE),
  n_miss_frlunch = sum(is.na(fr_lunch)),
  mean_fr_lunch=mean(fr_lunch, na.rm = TRUE))
```

Student exercise

summarise(): counts with logical vectors, part II

Application: count number in a group that satisfy some condition

Task: For each combination of event_type and event_inst, how many visits to communities that are majority Latinx or Black?

```
#event_berk %>% select(pct_black_zip, pct_hispanic_zip)
event_berk %>% group_by (event_inst, event_type) %>% summarise(
  n_events=n(), # number of events by group
  n_nonmiss_latbl = sum(!is.na(pct_black_zip) & !is.na(pct_hispanic_zip)), # w/
  n_majority_latbl= sum(pct_black_zip+ pct_hispanic_zip>50, na.rm = TRUE)) # num

#> # A tibble: 10 x 5
#> # Groups:   event_inst [?]
#>   event_inst event_type  n_events n_nonmiss_latbl n_majority_latbl
#>   <chr>      <fct>      <int>      <int>          <int>
#> 1 In-State  public hs          260          259            88
#> 2 In-State  private hs           36           36             7
#> 3 In-State  2yr college       107          102            27
#> 4 In-State  4yr college         12           10             0
#> 5 In-State  other                50           49            17
#> 6 Out-State public hs          184          184            27
#> 7 Out-State private hs       135          133            20
#> 8 Out-State 2yr college         1            1             0
#> 9 Out-State 4yr college         3            2             0
#> 10 Out-State other              90           87            20
```

summarise(): proportions with logical values

Application: count proportion of obs in group that satisfy some condition

- o Syntax: `group_by(vars) %>% summarise(prop = mean(TRUE/FALSE condition))`

Task:

- o separately for in-state/out-of-state, what proportion of visits to public high schools are to communities with median income greater than \$100,000?

Steps:

1. Filter public HS visits
2. group by in-state vs. out-of-state
3. Create measure

```
event_berk %>% filter(event_type == "public hs") %>% # filter public hs visits
  group_by (event_inst) %>% # group by in-state vs. out-of-state
  summarise(
    n_events=n(), # number of events by group
    n_nonmiss_inc = sum(!is.na(med_inc)), # w/ nonmissings values median inc,
    p_incgt100k = mean(med_inc>100000, na.rm=TRUE)) # proportion visits to $100K
#> # A tibble: 2 x 4
#>   event_inst n_events n_nonmiss_inc p_incgt100k
#>   <chr>      <int>      <int>      <dbl>
#> 1 In-State      260        257        0.272
#> 2 Out-State     184        184        0.516
```

What if we forgot to put `na.rm=TRUE`?

Student exercise

summarise(): Other functions

Common functions to use with `summarise`:

Function	Description
<code>n</code>	count
<code>n_distinct</code>	count unique values
<code>mean</code>	mean
<code>median</code>	median
<code>max</code>	largest value
<code>min</code>	smallest value
<code>sd</code>	standard deviation
<code>sum</code>	sum of values
<code>first</code>	first value
<code>last</code>	last value
<code>nth</code>	nth value
<code>any</code>	condition true for at least one value?

Note: These functions can also be used on their own or with `mutate()`

summarise(): Other functions

Maximum value in a group

```
max(c(10,50,8))  
#> [1] 50
```

Task: For each combination of in-state/out-of-state and event type, what is the maximum value of `med_inc`?

```
event_berk %>% group_by(event_type, event_inst) %>%  
  summarise(max_inc = max(med_inc))  
#> # A tibble: 10 x 3  
#> # Groups:   event_type [?]  
#>   event_type event_inst max_inc  
#>   <fct>      <chr>      <dbl>  
#> 1 public hs    In-State      NA  
#> 2 public hs    Out-State  223556.  
#> 3 private hs   In-State  250001  
#> 4 private hs   Out-State    NA  
#> 5 2yr college In-State      NA  
#> 6 2yr college Out-State  153070.  
#> 7 4yr college In-State      NA  
#> 8 4yr college Out-State      NA  
#> 9 other       In-State      NA  
#> 10 other      Out-State      NA
```

What did we do wrong here?

summarise(): Other functions

Isolate first/last/nth observation in a group

```
x <- c(10,15,20,25,30)
first(x)
last(x)
nth(x,1)
nth(x,3)
nth(x,10)
```

Task: after sorting event_berk by [arrange()] by event_type and event_datetime_start, what is the value of event_date for:

- first event for each event type?
- the last event for each event type?
- the 10th event for each event type?

```
event_berk %>% arrange(event_type, event_datetime_start) %>%
  group_by(event_type) %>%
  summarise(
    n_events = n(),
    date_first= first(event_date),
    date_last= last(event_date),
    date_50th= nth(event_date, 50)
  )
```

Student exercise

something that involves whether visits adhered to a certain pattern? e.g., visited org of type 1 and then org of type 2 in succession?

Attach aggregate measures to your data frame

We can attach aggregate measures to a data frame by using `group_by` without `summarise()`

Example task:

- Using `event_berk` data frame, create (1) a measure of average income across all events and (2) a measure of average income for each event type

Create measure of average income across all events

```
event_berk_temp <- event_berk %>%  
  arrange(event_date) %>% # sort by event_date (optional)  
  select(event_date, event_type, med_inc) %>% # select vars to be retained (optional)  
  mutate(avg_inc = mean(med_inc, na.rm=TRUE)) # create avg. inc measure  
  
dim(event_berk_temp)  
event_berk_temp %>% head(5)
```

Create measure of average income by event type

```
event_berk_temp <- event_berk_temp %>%  
  group_by(event_type) %>% # grouping by event type  
  mutate(avg_inc_type = mean(med_inc, na.rm=TRUE)) # create avg. inc measure  
  
event_berk_temp %>% head(5)
```

Attach aggregate measures to your data frame

Task: Create a measure that identifies whether `med_inc` associated with the event is higher/lower than average income for all events of that type

Steps:

1. Create measure of average income for each event type [already done]
2. Create measure that compares income to average income for event type

```
# average income at recruiting events across all universities
event_berk_tempv2 <- event_berk_temp %>%
  mutate(gt_avg_inc_type = med_inc > avg_inc_type) %>%
  select(-(avg_inc)) # drop avg_inc (optional)
event_berk_tempv2 # note how med_ic = NA are treated
```

create integer indicator rather than logical

```
event_berk_tempv2 <- event_berk_tempv2 %>%
  mutate(gt_avg_inc_type = as.integer(med_inc > avg_inc_type))
event_berk_tempv2 %>% head(4)
#> # A tibble: 4 x 5
#> # Groups:   event_type [3]
#>   event_date event_type med_inc avg_inc_type gt_avg_inc_type
#>   <date>      <fct>      <dbl>      <dbl>          <int>
#> 1 2017-04-11 other        62022.      70922.           0
#> 2 2017-04-14 public hs    125056.      93956.           1
#> 3 2017-04-23 private hs    78897       88695.           0
#> 4 2017-04-23 other         55127       70922.           0
```

Student exercise

Task: is `pct_white_zip` at a particular event higher or lower than the average `pct_white_zip` for that `event_type`?

- Note: all events attached to a particular `zip_code`
- `pct_white_zip`: pct of people in that `zip_code` who identify as white

Steps in task:

- Create measure of average `pct white` for each `event_type`
- Compare whether `pct_white_zip` is higher or lower than this average